Polarized User and Topic Tracking in Twitter

Mauro Coletto IMT Lucca, ISTI–CNR Pisa mauro.coletto@isti.cnr.it

Salvatore Orlando DAIS-University of Venice orlando@unive.it Claudio Lucchese ISTI-CNR Pisa claudio.lucchese@isti.cnr.it

Raffaele Perego ISTI-CNR Pisa raffaele.perego@isti.cnr.it

ABSTRACT

Digital traces of conversations in micro-blogging platforms and OSNs provide information about user opinion with a high degree of resolution. These information sources can be exploited to understand and monitor collective behaviors. In this work, we focus on polarization classes, i.e., those topics that require the user to side exclusively with one position. The proposed method provides an iterative classification of users and keywords: first, polarized users are identified, then polarized keywords are discovered by monitoring the activities of previously classified users. This method thus allows tracking users and topics over time. We report several experiments conducted on two Twitter datasets during political election time-frames. We measure the user classification accuracy on a golden set of users, and analyze the relevance of the extracted keywords for the ongoing political discussion.

1. INTRODUCTION

Recently, the analysis of blogging platforms and streaming information sources (e.g., Twitter) has received great attention in the Information Retrieval and in the Data Mining communities. We focus on the frequent scenario where users interact and produce contents according to a set of polarization classes. By polarization classes we mean subjects that require the user to side exclusively with one part. Political parties are typical examples of these classes. Other examples include brand analysis, products comparison, and opinion mining in general. In these scenarios the polarization classes are known, and some limited information may also be available, e.g., a set of relevant keywords. This limited knowledge allows us to restrict the scope of the analysis, but several challenging tasks are left open. The first is how to identify the users being polarized (or not) according to those classes. The second task concerns the identification of the most relevant sub-topics being discussed among such users. The third is how to monitor the evolution of such user communities and their on-line discussions over time. Those tasks are all very challenging as the available knowledge may be approximate or insufficient, and it may also become obsolete over time. Therefore, the classification into polarization classes should be able to self-update continuously by catching upcoming relevant

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SIGIR '16, July 17 - 21, 2016, Pisa, Italy

© 2016 Copyright held by the owner/author(s). Publication rights licensed to ACM. ISBN 978-1-4503-4069-4/16/07...\$15.00

DOI: http://dx.doi.org/10.1145/2911451.2914716

users and discussion topics. The present work is related to the Topic Detection and Tracking (TDT) subject [2], which has been widely explored within the scope of news stream analysis [11]. We focus on content and user tracking for polarized users. This notion is connected with the concept of controversy in Social Media, which have been studied, mostly in political contexts, using data coming from different sources [1, 4, 6, 8]. Another related research area is trending topics analysis. Various trend detection models are proposed in [7,9]. Our approach is different in several regards from current literature, since we rather focus on the identification of polarized communities. In our experiments we use electoral data from Twitter. In this case, the polarization classes are political parties and candidates. Several works analyzed the opportunities and limitations in using Twitter as a predictor of an election's outcome [3, 5, 10]. Our goal is completely different, as we do not draw any conclusion about the expected share of votes for the given parties or candidates. We use this specific kind of data, as it is a typical example of polarized users. We show that the proposed algorithm is able to identify polarized users, by also analyzing the ongoing discussions among the respective communities.

The main contribution of this work is a new iterative algorithm, named PTR (Polarization TRacker), for the discovery of *polarized users* in a Twitter stream, and a temporal version TPTR (Temporal PTR), able to track users and topics over time. While there exist several works about community detection and trending topic tracking, we propose a novel setting where the number of communities is known, but very little information is provided (a keyword per class only), and those communities are competing with each other. We conduct an objective evaluation of the proposed algorithms by measuring their classification accuracy on a *golden* set of users.

2. USER AND TOPIC TRACKING

Let $\mathcal{T}=\{t_1,t_2,\ldots\}$ be the stream of tweets generated by the set of users $\mathcal{U}=\{u_1,u_2,\ldots\}$. We focus on the analysis of user behavior with respect to a set of *polarization classes* \mathcal{C} . The goal of this work is thus to build a *partitional clustering* of the Twitter users, where each of the clusters is associated by construction with a single polarization class (or unassigned). Our method can be seen as a *semi-supervised clustering* one, although, unlike classic methods, we do not provide any class representative around which the final clustering is induced. Indeed, the proposed method is only *loosely* supervised as the only knowledge available is the number of classes, and a short class description (a keyword).

An important issue is the evaluation of our algorithm. To this end, we exploit a *golden set* of polarized users, each unequivocally associated with a class $c \in \mathcal{C}$. Note that such knowledge is not exploited to train a classifier, but only for evaluation purpose.

Algorithm 1 User Classification Algorithm

```
Require: The set of polarized hashtags H_c and the previously found
                set of polarized users U_c^* for each class c \in \mathcal{C}
Ensure: New set of polarized users \{U_c\}_{c \in \mathcal{C}}
1: procedure USERSCLASS( \{H_c\}_{c \in \mathcal{C}}, \{U_c^*\}_{c \in \mathcal{C}} )
2:
           for u \in \mathcal{U}, c \in \mathcal{C} do
                                                                                             ⊳ Find polarized tweets
                 T_{u,c} = \{ t \in \mathcal{T}_u \mid \mathcal{H}_t \cap H_c \neq \emptyset \land \mathcal{H}_t \cap H_{c' \neq c} = \emptyset \}
4:
5:
           for c \in \mathcal{C} do
                 U_c \leftarrow \emptyset
6:
7:
           for u \in \mathcal{U} do
                                                                                       ▷ Check user's polarization
                 if \exists c \in \mathcal{C} \mid \forall c' \in \mathcal{C}, c' \neq c \quad |T_{u,c}| > \alpha \cdot |T_{u,c'}| then
8:
                       U_c \leftarrow U_c \cup u
9:
                 else if \exists c \in \mathcal{C} \mid u \in U_c^* then
10:
                        U_c \leftarrow U_c \cup u
             return \{U_c\}_{c\in\mathcal{C}}
11:
```

2.1 The PTR algorithm

The Polarization TRacker (PTR) algorithm requires some initial *seed topics* that identify the classes of interests. We propose to identify them with a single textual keyword for each class $c \in \mathcal{C}$. Although each keyword identifies a topic, e.g., a political party, it is not sufficient to correctly classify users, as all these seed topics are likely to be mentioned in many users' tweets, e.g., to contrast the achievements of a given party with the deficiencies of the others. Without loss of generality, we limit our keyword selection to Twitter *hashtags*. Therefore, the single textual keyword we initially choose for each class c is a single hashtag appearing in the user tweets, and around them we start identifying the user clusters. The final goal is to extract the best discriminating hashtags that are able to identify the actual clusters of polarized users, who belong with high probability to one of the classes $c \in \mathcal{C}$.

We denote the representative hashtags, one for each $c \in \mathcal{C}$, called seed hashtags, by $H_c^{\tau=0}$, where τ is the algorithm's iteration number. Note that each initial set $H_c^{\tau=0}$, one for each c, is not necessarily composed of a discriminating hashtag. This set $H_c^{\tau=0}$ is then used to classify polarized users on the basis of their use of the seed hashtags. We denote by $U_c^{\tau+1}$ the clusters of users in $\mathcal U$ that are identified as belonging to class c, according to their tweets and to the given hashtags H_c^{τ} . Similarly, the new hashtags $H_c^{\tau+1}$ are generated by finding those that best discriminate the users in $U_c^{\tau+1}$. This refinement process is iterated for all $c \in \mathcal C$: from hashtags $\{H_c^{\tau}\}_{c\in\mathcal C}$ to users $\{U_c^{\tau+1}\}_{c\in\mathcal C}$, and finally to hashtags $\{H_c^{\tau}\}_{c\in\mathcal C}$. The algorithm terminates when H_c^{τ} converges.

Specifically, PTR iterates the two classification steps USERCLASS and HASHTAGSCLASS. Algorithm 1 illustrates the former step of the iterative process¹. The goal of this step is to identify polarized users on the basis of the given hashtags. First, we identify polarized tweets, which mention hashtags in H_c . We consider the classification of each single tweet t by considering all the mentioned hashtags \mathcal{H}_t , as we believe each tweet is a very relevant expression of a user's thought on a specific topic. Since we are interested in polarized users, with the goal of achieving high precision we discard all the tweets which contain hashtags belonging to more than one set $\{H_c\}_{c\in\mathcal{C}}$. For each user $u\in\mathcal{U}$ and for each class $c\in\mathcal{C}$ we denote the set of polarized tweets by $T_{u,c}$. We thus measure the user polarization: if for some classes c, the number of tweets in $T_{u,c}$ is significantly larger than for any other class (parameter α), then the user is labeled with the class c and added to the set of polarized users U_c (see line 7). Note that the user classification is intended to be an update of the classification conducted during the previous step. The goal the second step is to process all the hashtags adopted by classified users U_c in order to discover a new set of

Algorithm 2 Hashtag Classification Algorithm

```
 \begin{array}{lll} \textbf{Require:} & \text{The set of polarized users } U_c \text{ for each class } c \in \mathcal{C} \\ \textbf{Ensure:} & \text{Polarized hashtags } H_c \\ 1: & \textbf{procedure HASHTAGSCLASS}(& \{U_c\}_{c \in \mathcal{C}}) \\ 2: & \textbf{for } c \in \mathcal{C} \textbf{ do} \\ 3: & H_c \leftarrow \emptyset \\ 4: & H_c^* \leftarrow \bigcup_{u \in U_c} \mathcal{H}_u \\ 5: & \textbf{for } h \in \bigcup_{c \in \mathcal{C}} H_c^* \textbf{ do} \\ 6: & \textbf{if } \exists c \mid \forall c' \neq c \quad S_c(h) > \beta \cdot S_{c'}(h) \textbf{ then} \\ 7: & H_c \leftarrow H_c \cup h \\ 8: & \textbf{return } \{H_c\}_{c \in \mathcal{C}} \\ \end{array}
```

discriminating hashtags H_c , as illustrated in Alg. 2. In order to detect $\{H_c\}_{c\in\mathcal{C}}$, we take into considerations all the hashtags \mathcal{H}_u used by any user $u \in U_c$, and not only those occurring in the polarized tweets $T_{u,c}$ (line 4). This allows to extend our analysis to the full set of topics discussed by the users, even if they were not captured in the early iterations of the algorithm. First, for each $c \in C$ we retrieve the set of hashtags used by the users in U_c , considering all their tweets, denoted by \mathcal{T}_c , independent of the classification of the single tweets in the previous iteration. In our experiments we consider the top frequent 500 hashtags in \mathcal{T}_c . Given the resulting set of candidate hashtags for each $c \in \mathcal{C}$, namely H_c^* , we extract from them the new hashtags that highly discriminate each class c, and these are eventually added to the new set H_c (line 7). Specifically, the discriminating hashtags are those highly used by the current set of users U_c , and partially used by any other user in $U_{c'}$, $c' \neq c$. We define a function $S_c(h)$ to measure the goodness of hashtag h for each community of polarized users U_c . Let \mathcal{T}_h be the set of tweets in \mathcal{T} mentioning hashtag h, independent of the users who posted these tweets. Moreover, let $\mathcal{T}_{H_o^*}$ be the set of tweets in \mathcal{T} containing at least one hashtag in the set H_c^* . We score the goodness of a hashtag for a polarization class as follows:

$$S_c(h) = \frac{|\mathcal{T}_h \cap \mathcal{T}_{H_c^*}|}{|\mathcal{T}_{H_c^*}|} \cdot \prod_{c' \in \mathcal{C}, c' \neq c} \left(1 - \frac{|\mathcal{T}_h \cap \mathcal{T}_{H_{c'}^*}|}{|\mathcal{T}_{H_{c'}^*}|}\right)$$

where we consider the naive hypothesis of independent occurrence of the hashtags in the various sets. In practice, $S_c(h)$ is the probability of seeing h only in H_c^* , whereas h is not present in all the other sets of hashtags $H_{c'\neq c}^*$. Given a hashtag h, the score $S_c(h)$ is used to rank the various classes, thus assigning h to class with the highest score. Since we aim at promoting high discriminating hashtags, not only we assign the hashtag h having the highest $S_c(h)$ to the new set H_c , but only if $S_c(h) > \beta \cdot S_{c'}(h)$, $\forall c' \neq c$, where $\beta \geq 1$. Note that if a tie exists between the to 2-top scores classes, the hashtag h is not assigned to any H_c , since it is considered not discriminating enough.

3. EXPERIMENTAL EVALUATION

3.1 Data collection and cleansing

We use two Twitter datasets related to political elections that recently took place in Italy. **Dataset IT13**: data about primary election for largest social democratic political party in Italy (PD), which took place in December 2013 with 3 candidates: Mr. Renzi, Mr. Cuperlo, and Mr. Civati. **Dataset EU14**: data about European Parliament election held in Italy in May 2014². The data are collected through Twitter API by querying a list of keywords related to the topic and the candidates, large enough to guarantee a good coverage

 $^{^{1}}$ Note that we omitted the superscript au for the sake of simplifying the notation.

²The main national parties connected to different European political groups were: *Partito Democratico* (PD), *Movimento 5 Stelle* (M5S), *Forza Italia* (FI), *Lega Nord* (LN), *Tsipras* (AET). We ignore smaller parties and NCD-UDC for its limited presence in Twitter.

Table 1: Data Statistics

(a) Full dataset

Dataset	IT13	EU14
tweets in original raw data	1.7 M.	2.3 M.
pre-electoral tweets ${\mathcal T}$	95,627	364,132
users with $ \mathcal{H}_u > 0$	11,368 (65%)	28,340 (56%)

(b) Golden dataset

Da	taset IT13	3	Dataset EU14			
\mathcal{C}	Tweets	Users	C Tweets User			
Renzi	330	109	PD	262	129	
Cuperlo	4759	243	M5S	146	95	
Civati	2925	700	FI	1263	199	
			LN	480	226	
			AET	757	328	
total	8014	1052	total	2908	977	

of the elections. Both final datasets cover 9 days before the election day. We discard partial data and potentially irrelevant tweets, considering only tweets being geo-located and in Italian language. Table 1b reports some information about the two datasets.

3.2 Evaluation of the PTR algorithm

We build an evaluation dataset by identifying those users whose opinion can be inferred with high confidence. During elections, as for other events, very specific hashtags are used over Twitter to express a strong intention of vote or an explicit membership in a group. We assume that users that frequently use one of such hashtags are strongly sided with one of the competing parties and they will not change idea in the short term. Such hashtags, named golden hashtags, are handpicked among the 500 most frequent in the data. The used golden hashtags are of the kind #IVoteParty. We identify one/two golden hashtags per class $c \in \mathcal{C}$ both in the EU14 (e.g. #IVoteTsipras for AET) and in the IT13 (e.g. #prefeRenzi for Renzi) dataset. The set of reference users were identified by applying Algorithm 1 with the above golden hashtags as input. This guarantees that a user is safely considered as polarized to a party $c \in \mathcal{C}$ if her tweets contain only one of the golden hashtags associated with the various classes $c \in \mathcal{C}$. We denote with $\mathcal{Z} = \{z_1, z_2, \ldots\}$ this set of polarized users, and with $Z_c \subseteq \mathcal{Z}$ those supporting a specific formation c (Z_c is a partitioning of \mathcal{Z}). The composition of resulting golden dataset is reported in Table 1a. The golden dataset is thus a small fraction of the full dataset. A global analysis of the Twitter stream cannot be based on a few very polarized hashtags. Note that the relative popularity of the parties is not simply proportional to the number of votes received, but it depends on the efficacy of the hashtag promoted. We remark that, for the sake of fairness, we remove the golden hashtags from the datasets before the application of any algorithm. The set of users \mathcal{Z} in the golden dataset, is used to evaluate the users classification accuracy of the proposed method. Given the users classification U_c provided by some given algorithm, precision, recall and F-Measure are restricted to the set \mathcal{Z} . Formally, for any given class $c \in \mathcal{C}$, precision and recall are defined as:

$$P_c(U_c) = \frac{|U_c \cap Z_c|}{|U_c \cap Z|}$$
 $R_c(U_c) = \frac{|U_c \cap Z_c|}{|Z_c|}$

The F-measure F_c is the harmonic means of P_c and R_c . The macro F-measure average over the classes $c \in \mathcal{C}$ is denoted with F. In addition, as the proposed algorithm may not be able to classify all of the users in \mathcal{Z} , we report also the user coverage γ and Γ on both

the golden set and the overall dataset respectively:

$$\gamma(U = \cup_{c \in \mathcal{C}} U_c) = \frac{|U \cap \mathcal{Z}|}{|\mathcal{Z}|} \qquad \Gamma(U = \cup_{c \in \mathcal{C}} U_c) = \frac{|U|}{|\mathcal{U}|}$$

As a baseline we use the k-means clustering algorithm. Each user u is represented by a vector of 500 features, corresponding to the 500 most frequent hashtags in the dataset. The user feature vector stores the frequency of a hashtag in the stream of tweets \mathcal{T}_u published by the user. We discard users who do not use any hashtag in their tweets. We normalize the feature vectors for each user to unit L^2 norm. We impose the number of the clusters k equal to the number of classes $|\mathcal{C}|$ and, to simulate the same starting condition of our method, we built the initial centroids so as to encode the seed hashtags. The centroid for a class c is thus a vector with a single 1 in the position of the seed hashtag, and 0 otherwise. The result of the k-means baseline is thus a clustering of users based on the seed hashtags provided. Table 2a reports the results of the k-means baseline. F-measure values are low for the IT13 dataset. For instance, k-means provides low accuracy and recall for the first class. This is mainly due to the fact that the hashtags corresponding to popular parties or candidates are very often used by different users, regardless of their orientation. In other cases (e.g., LN and AET), the hashtags are used mostly within the respective communities.

In the following, we analyze in detail the iteration-by-iteration behavior of the proposed PTR algorithm. We test our algorithm by setting $\alpha = 2$ and $\beta = 1$, after a tuning step. During the first iteration, PTR is fed with the seed hashtags. Algorithm 1 uses those hashtags to find a subset of polarized users in \mathcal{U} . This step is similar to other works, where mentions of a party or candidate are used to estimate their popularity or to classify users [3, 10]. Unlike other approaches, PTR aims at discovering a subset of polarized users, thus requiring, that a user mentions a party at least twice any other. The results of such user classification are evaluated over the golden dataset, as reported in the first line of Table 3a. Regarding average precision, PTR is already significantly superior to the k-means baseline for IT13 dataset. This is already surprising, as the seed hashtags are very generic. On the other hand, the k-means baseline might be negatively affected by the sparsity of the data. The results are different on the two datasets in terms of average recall. PTR has similar performance to k-means on the IT13 dataset, while the recall is significantly lower on the EU14 dataset. This is

Table 2: Comparison with the Baseline.

(a) k-means baseline performance

Dataset IT13				Dataset EU14			
\mathcal{C}	P_c	R_c	F_c	\mathcal{C}	P_c	R_c	F_c
Renzi	0.144	0.257	0.185	PD	0.536	0.457	0.493
Cuperlo	0.252	0.543	0.344	M5S	0.359	0.895	0.512
Civati	0.766	0.366	0.495	FI	0.495	0.734	0.591
				LN	0.995	0.916	0.954
				AET	1.000	0.387	0.558
avg.	0.387	0.389	0.341	avg.	0.677	0.678	0.622
$\gamma = 1.0$ $\Gamma = 0.653$			γ	= 1.0	$\Gamma =$	0.557	

(b) PTR Iteration-2 performance

Dataset IT13				Dataset EU14			
С	P_c	R_c	F_c	\mathcal{C}	P_c	R_c	F_c
Renzi	0.350	0.752	0.478	PD	0.733	0.488	0.586
Cuperlo	0.869	0.300	0.446	M5S	0.325	0.842	0.469
Civati	0.916	0.747	0.823	FI	0.955	0.533	0.684
				LN	0.981	0.938	0.959
				AET	0.974	0.451	0.617
avg.	0.712	0.600	0.582	avg.	0.794	0.650	0.663
$\gamma =$	0.845	Γ =	= 0.532	$\gamma =$	0.830	$\Gamma =$	0.367

Table 3: Algorithm Performance.

(a) PTR iteration by iteration performance

	Dataset IT13			Dataset EU14		
Iter	F	γ	Γ	F	γ	Γ
1	0.358	0.490	0.218	0.514	0.670	0.163
2	0.582	0.845	0.522	0.663	0.830	0.367
3	0.588	0.853	0.532	0.662	0.831	0.386
4	0.588	0.853	0.534	0.661	0.834	0.390

(b) TPTR day-by-day performance

	Dataset IT13			Dataset EU14		
Day	F	γ	Γ	F	γ	Γ
1	0.177	0.199	0.045	0.155	0.164	0.025
2	0.225	0.348	0.114	0.464	0.465	0.079
3	0.304	0.457	0.166	0.529	0.570	0.116
4	0.333	0.563	0.234	0.585	0.671	0.180
5	0.368	0.606	0.261	0.588	0.726	0.235
6	0.397	0.671	0.315	0.574	0.762	0.269
7	0.387	0.721	0.363	0.596	0.794	0.302
8	0.387	0.765	0.408	0.637	0.846	0.334
9	0.391	0.811	0.461	0.635	0.876	0.349

confirmed by the coverage values γ and Γ . In comparison with the baseline, the performance of PTR in terms of macro F-measure is satisfactory on the IT13 dataset, but not on the EU14 dataset yet. The output of the first iteration is a new set of hashtags which is exploited in the next iteration. By looking at the best scoring hashtag, we can already observe an interesting behavior of the algorithm for some $c \in \mathcal{C}$. In dataset EU14, the best tags for FI and LN are the leaders of the respective parties, detecting that the original seed hashtags are not discriminating in this case. In Table 2b we report in detail the results after the second iteration of PTR. The first interesting result is that the average recall is significantly higher on both datasets. This is due to the new hashtags discovered in addition to the *seed* ones during the previous iteration, which, in turn, lead to the identification of a larger set of users: the coverage γ is now beyond 80% on the *golden* set, and Γ has doubled in this iteration. Also the average precision is higher w.r.t. the previous iteration scoring more than 0.7. This is both because of the increased number of classified users, and of the updated user classification. As a result, the F-measure has an overall improvement w.r.t. the k-means baseline of +71% and +7% on datasets IT13 and EU14 respectively. As shown in Table 3a PTR becomes stable very early. The largest improvement is achieved with the second iterations. This means that the most relevant hashtags are discovered early, and only slight changes occur afterwards. The subsequent iterations marginally increase the number of classified users. Note that the algorithm is classifying the polarized users found in the whole set U. PTR found about 6.7 and 27 thousands polarized users on the dataset IT13 and EU14 respectively. We conclude that in most cases, two iterations of the algorithm provide sufficient classification quality. For the lack of space we can not report a exhaustive qualitative analysis of the outcome, but we observe that the procedure is able to extract relevant keywords: namely prominent politicians, the party itself and political mottoes characterizing each c in the political scene.

We finally propose a variant of PTR, that is TPTR (temporal PTR), to perform the tracking of topics and users in time. In our case we consider the evolution day by day. The procedure follows Algorithm 1 and Algorithm 2 with the difference that at iteration τ only the tweets \mathcal{T}_u written in the τ -th day are considered. We perform TPTR on IT13 and on EU14 datasets. In Table 3b the evaluation of the temporal iterative procedure is shown. The macro

F-measure is increasing day by day both for the effect of a better classification and for the presence of new users. Note that we evaluate the time iterative method day by day on the entire *golden set* of users. F-measure values are low because not all users in the *golden set* were active every day.

4. CONCLUSION

We propose a novel algorithm for the simultaneous tracking of polarized users and discriminating topics in OSNs. Specifically, it iteratively detects polarized users, and from their contents the discussed discriminating topics. We also introduce a temporal variant, where the information extracted during one day of analysis is exploited for the next day. Indeed, the classification of users makes the algorithm more robust in terms of concept drifts, as new trends may be detected as early as they pop up. At the same time, the identification of discriminating topics helps in detecting users moving from one class to another. The algorithm is tested on two Twitter data samples. We evaluate the quality of user classification on a golden set of users, showing significant improvements over the baseline. The proposed methodology is general and it can be applied to different scenarios. We believe that this methodology based on polarization may also impact on broad area of social network analysis, e.g., by complementing the proposed classification with community detection and information diffusion over time. As a future work, we aim to improve the temporal analysis dealing with streaming data.

Supported by EC H2020 Program INFRAIA-1-2014-2015 (654024).

5. REFERENCES

- [1] L. A. Adamic and N. Glance. The political blogosphere and the 2004 us election: divided they blog. In *Proc. of the 3rd Workshop on Link discovery*, pages 36–43. ACM, 2005.
- [2] J. Allan. Topic detection and tracking: event-based information organization, volume 12. Springer Science & Business Media, 2012.
- [3] M. Coletto, C. Lucchese, S. Orlando, and R. Perego. Electoral predictions with twitter: a machine-learning approach. *IIR*, 2015.
- [4] K. Garimella, G. De Francisci Morales, A. Gionis, and M. Mathioudakis. Quantifying controversy in social media. In ACM WSDM, pages 33–42, 2016.
- [5] D. Gayo-Avello, P. T. Metaxas, and E. Mustafaraj. Limits of electoral predictions using twitter. In ICWSM, 2011.
- [6] C. V. Gysel, B. Goethals, and M. de Rijke. Determining the presence of political parties in social circles. In AAAI ICWSM, 2015.
- [7] J. Lin, R. Snow, and W. Morgan. Smoothing techniques for adaptive online language models: topic tracking in tweet streams. In ACM SIGKDD, pages 422–429. ACM, 2011.
- [8] A. Makazhanov, D. Rafiei, and M. Waqar. Predicting political preference of twitter users. *Social Network Analysis* and Mining, 4(1):1–15, 2014.
- [9] M. Mathioudakis and N. Koudas. Twittermonitor: trend detection over the twitter stream. In *Proc. of the 2010 ACM SIGMOD*, pages 1155–1158. ACM, 2010.
- [10] A. Tumasjan, T. O. Sprenger, P. G. Sandner, and I. M. Welpe. Predicting elections with twitter: What 140 characters reveal about political sentiment. *ICWSM*, 10:178–185, 2010.
- [11] F. Walls, H. Jin, S. Sista, and R. Schwartz. Topic detection in broadcast news. In *Proceedings of the DARPA broadcast* news workshop, pages 193–198, 1999.